

Lecture 13: Dijkstra's Algorithm

Review

- Single-Source Shortest Paths on weighted graphs
- Previously: $O(|V| + |E|)$ -time algorithms for small positive weights or DAGs
- Last time: Bellman-Ford, $O(|V||E|)$ -time algorithm for **general graphs** with **negative weights**
- Today: faster for **general graphs** with **non-negative edge weights**, i.e., for $e \in E$, $w(e) \geq 0$

Restrictions		SSSP Algorithm		
Graph	Weights	Name	Running Time $O(\cdot)$	Lecture
General	Unweighted	BFS	$ V + E $	L09
DAG	Any	DAG Relaxation	$ V + E $	L11
General	Any	Bellman-Ford	$ V \cdot E $	L12
General	Non-negative	Dijkstra	$ V \log V + E $	L13 (Today!)

Non-negative Edge Weights

- **Idea!** Generalize BFS approach to weighted graphs:
 - Grow a sphere centered at source s
 - Repeatedly explore closer vertices before further ones
 - But how to explore closer vertices if you don't know distances beforehand? : (
- **Observation 1:** If weights non-negative, monotonic distance increase along shortest paths
 - i.e., if vertex u appears on a shortest path from s to v , then $\delta(s, u) \leq \delta(s, v)$
 - Let $V_x \subset V$ be the subset of vertices reachable within distance $\leq x$ from s
 - If $v \in V_x$, then any shortest path from s to v only contains vertices from V_x
 - Perhaps grow V_x one vertex at a time! (but growing for every x is slow if weights large)
- **Observation 2:** Can solve SSSP fast if given order of vertices in increasing distance from s
 - Remove edges that go against this order (since cannot participate in shortest paths)
 - May still have cycles if zero-weight edges: repeatedly collapse into single vertices
 - Compute $\delta(s, v)$ for each $v \in V$ using DAG relaxation in $O(|V| + |E|)$ time

Dijkstra's Algorithm

- Named for famous Dutch computer scientist **Edsger Dijkstra** (actually Dijkstra!)

11 August 1982
 prof. dr. Edsger W. Dijkstra
 Burroughs Research Fellow

- **Idea!** Relax edges from each vertex in increasing order of distance from source s
- **Idea!** Efficiently find next vertex in the order using a data structure
- **Changeable Priority Queue** Q on items with keys and unique IDs, supporting operations:

<code>Q.build(X)</code>	initialize Q with items in iterator x
<code>Q.delete_min()</code>	remove an item with minimum key
<code>Q.decrease_key(id, k)</code>	find stored item with ID id and change key to k


- Implement by **cross-linking** a Priority Queue Q' and a Dictionary D mapping IDs into Q'
- Assume vertex IDs are integers from 0 to $|V| - 1$ so can use a direct access array for D
- For brevity, say item x is the tuple $(x.id, x.key)$

- Set $d(s, v) = \infty$ for all $v \in V$, then set $d(s, s) = 0$
- Build changeable priority queue Q with an item $(v, d(s, v))$ for each vertex $v \in V$
- While Q not empty, delete an item $(u, d(s, u))$ from Q that has minimum $d(s, u)$
 - For vertex v in outgoing adjacencies $\text{Adj}^+(u)$:
 - * If $d(s, v) > d(s, u) + w(u, v)$:
 - Relax edge (u, v) , i.e., set $d(s, v) = d(s, u) + w(u, v)$
 - Decrease the key of v in Q to new estimate $d(s, v)$

- Run Dijkstra on example

Example

Delete v from Q	$d(s, v)$				
	s	a	b	c	d
s	0	∞	∞	∞	∞
c		10	∞	3	∞
d		7	11		5
a		7	10		
b			9		
$\delta(s, v)$	0	7	9	3	5



Correctness

- **Claim:** At end of Dijkstra's algorithm, $d(s, v) = \delta(s, v)$ for all $v \in V$
- **Proof:**
 - If relaxation sets $d(s, v)$ to $\delta(s, v)$, then $d(s, v) = \delta(s, v)$ at the end of the algorithm
 - * Relaxation can only decrease estimates $d(s, v)$
 - * Relaxation is safe, i.e., maintains that each $d(s, v)$ is weight of a path to v (or ∞)
 - Suffices to show $d(s, v) = \delta(s, v)$ when vertex v is removed from Q
 - * Proof by induction on first k vertices removed from Q
 - * Base Case ($k = 1$): s is first vertex removed from Q , and $d(s, s) = 0 = \delta(s, s)$
 - * Inductive Step: Assume true for $k < k'$, consider k' th vertex v' removed from Q
 - * Consider some shortest path π from s to v' , with $w(\pi) = \delta(s, v')$
 - * Let (x, y) be the first edge in π where y is not among first $k' - 1$ (perhaps $y = v'$)
 - * When x was removed from Q , $d(s, x) = \delta(s, x)$ by induction, so:

$$\begin{aligned}
 d(s, y) &\leq \delta(s, x) + w(x, y) && \text{relaxed edge } (x, y) \text{ when removed } x \\
 &= \delta(s, y) && \text{subpaths of shortest paths are shortest paths} \\
 &\leq \delta(s, v') && \text{non-negative edge weights} \\
 &\leq d(s, v') && \text{relaxation is safe} \\
 &\leq d(s, y) && v' \text{ is vertex with minimum } d(s, v') \text{ in } Q
 \end{aligned}$$

- * So $d(s, v') = \delta(s, v')$, as desired

□

Running Time

- Count operations on changeable priority queue Q , assuming it contains n items:

Operation	Time	Occurrences in Dijkstra
$Q.\text{build}(X)$ ($n = X $)	B_n	1
$Q.\text{delete_min}()$	M_n	$ V $
$Q.\text{decrease_key}(id, k)$	D_n	$ E $

- Total running time is $O(B_{|V|} + |V| \cdot M_{|V|} + |E| \cdot D_{|V|})$
- Assume pruned graph to search only vertices reachable from the source, so $|V| = O(|E|)$

Priority Queue Q' on n items	Q Operations $O(\cdot)$			Dijkstra $O(\cdot)$ $n = V = O(E)$
	$\text{build}(X)$	$\text{delete_min}()$	$\text{decrease_key}(id, k)$	
Array	n	n	1	$ V ^2$
Binary Heap	n	$\log n_{(a)}$	$\log n$	$ E \log V $
Fibonacci Heap	n	$\log n_{(a)}$	$1_{(a)}$	$ E + V \log V $

- If graph is **dense**, i.e., $|E| = \Theta(|V|^2)$, using an Array for Q' yields $O(|V|^2)$ time
- If graph is **sparse**, i.e., $|E| = \Theta(|V|)$, using a Binary Heap for Q' yields $O(|V| \log |V|)$ time
- A Fibonacci Heap is theoretically good in all cases, but is not used much in practice
- We won't discuss Fibonacci Heaps in 6.006 (see 6.854 or CLRS chapter 19 for details)
- You should assume Dijkstra runs in $O(|E| + |V| \log |V|)$ time when using in theory problems

Summary: Weighted Single-Source Shortest Paths

Restrictions		SSSP Algorithm	
Graph	Weights	Name	Running Time $O(\cdot)$
General	Unweighted	BFS	$ V + E $
DAG	Any	DAG Relaxation	$ V + E $
General	Non-negative	Dijkstra	$ V \log V + E $
General	Any	Bellman-Ford	$ V \cdot E $

- What about All-Pairs Shortest Paths?
- Doing a SSSP algorithm $|V|$ times is actually pretty good, since output has size $O(|V|^2)$
- Can do better than $|V| \cdot O(|V| \cdot |E|)$ for general graphs with negative weights (next time!)